

Hetvi Shah

AI / Generative AI Engineer

Ahmedabad, Gujarat, India | +91 87580 33238 | hetvi4551@gmail.com | LinkedIn

PROFESSIONAL SUMMARY

AI / Generative AI Engineer with nearly 2 years of experience building production-grade **agentic AI systems, multi-agent workflows, RAG pipelines, and multilingual voice AI**. Hands-on experience with LangGraph, FastAPI, LiteLLM, Qdrant, model routing, guardrails, async processing, and Java/Python backend integration.

TECHNICAL SKILLS

AI/GenAI: LangGraph, LangChain, RAG, Agentic AI, Prompt Engineering, Tool Calling, Guardrails, Structured Outputs
LLM/Vector: LiteLLM, OpenAI, Gemini, Anthropic, DeepSeek, Qdrant, Hugging Face Embeddings
Voice AI: Pipecat, Deepgram, ElevenLabs, Twilio
Backend/Data: Python, FastAPI, Pydantic, asyncio, Java 17, Spring Boot, PostgreSQL, MySQL, Redis
Tools: n8n, Docker, Git, pytest, Postman, Swagger/OpenAPI, AWS S3

PROFESSIONAL EXPERIENCE

Wishtree Technologies

Associate Software Engineer

Software Engineer Intern

Ahmedabad, Gujarat

May 2025 – Present

Dec 2024 – Apr 2025

- Built client-facing GenAI solutions spanning **multi-agent systems, RAG, voice AI, and LLM-powered SaaS workflows** using LangGraph, FastAPI, vector search, and external tools/APIs.
- Designed reliable AI execution with checkpointing, retry/resume, model fallback, concurrency control, RPM/TPM rate limiting, Pydantic schemas, and PII/privacy validation.
- Integrated AI services with Spring Boot and Python backends, databases, authentication, third-party APIs, storage, and asynchronous workflows.

Zuru Tech

Frontend Developer Intern

Developed responsive web interfaces using HTML, CSS, and JavaScript.

Ahmedabad, Gujarat

May 2023 – Jun 2023

PROJECTS

Multiverz – Strategy Navigator

LangGraph, FastAPI, LiteLLM, Redis

AI-powered strategy platform using specialized agents to generate, evaluate, consolidate, and transform strategic insights into structured reports.

- Designed and built the production-grade LangGraph pipeline from scratch, migrating a complex n8n workflow to a stateful multi-agent architecture.
- Implemented checkpointing, retry/resume, step recovery, priority-based execution, concurrency control, and RPM/TPM rate limiting.
- Built resilient multi-model access using self-hosted LiteLLM, routing, fallback, cooldowns, and Tenacity-based exponential-backoff retries.

CYDRA Social

FastAPI, PostgreSQL, React, LLMs

AI-assisted social publishing platform for content generation, scheduling, and multi-platform publishing.

- Built AI-assisted content workflows and FastAPI backend services with PostgreSQL, JWT/OAuth, and LLM integrations.
- Implemented timezone-aware publishing to LinkedIn, X, and Meta, plus end-to-end Stripe subscription lifecycle workflows.

Multilingual Voice AI Agents

Pipecat, Deepgram, ElevenLabs, Twilio

Tool-enabled conversational AI POCs for healthcare, finance, travel, insurance, and commerce.

- Built inbound/outbound agents across 6 languages with real-time speech, LLM reasoning, telephony, and tool/API execution.
- Integrated workflows for appointments, loans, recovery, insurance claims, travel data, commerce, and automated emails.

AI Resource Finder

RAG, Qdrant, Hugging Face, Gemini

Natural-language employee search using semantic retrieval over skills and experience.

- Built the end-to-end RAG pipeline covering ingestion, chunking, embeddings, Qdrant indexing, semantic retrieval, and grounded Gemini responses.

EDUCATION

Dharmsinh Desai University – B.Tech in Information Technology

2021 – 2025

CGPA: **8.89/10** | AWS Academy Graduate – Cloud Foundations