

Hetvi, GenAI Engineer

Hetvi is a GenAI Engineer with hands-on experience building production-oriented AI solutions, voice Agents, and workflow automation using Python, FastAPI, LangGraph, n8n, and modern LLM APIs. She works on multi-agent systems, RAG, model routing, checkpointing, retries, REST APIs, webhooks, and system integration. She integrates AI capabilities across frontend and backend applications to streamline workflows and reduce manual effort. She focuses on building scalable, reliable, and maintainable AI solutions with strong attention to workflow resilience, observability, and production readiness.

SKILLS

- Designs and develops **GenAI applications, multi-agent systems, Voice AI, and workflow automation** using **Python, FastAPI, LangGraph, LangChain, n8n, and modern LLM APIs**, integrating AI capabilities into real business processes.
- Builds reliable AI workflows with **persistent state, structured outputs, checkpointing, retries, human-in-the-loop review, tool calling, model routing, model fallback handling, concurrency control, and rate limiting**.
- Integrates **AI solutions across frontend and backend applications** using **REST APIs, webhooks, databases, external platforms, and asynchronous processing** to support end-to-end AI workflows.
- Develops **RAG pipelines** using document ingestion, chunking, embeddings, **Qdrant/FAISS**, metadata filtering, and caching for context-aware and grounded AI responses.
- Implements **guardrails, validation, authentication, role-based access control, observability, logging, token tracking, Grafana, LangSmith, and Langfuse** to improve reliability, monitoring, and controlled AI usage.
- Works with **Python, Java, Spring Boot, React, Docker, AWS, Redis, Git, and SonarQube** to build scalable applications, integrate AI-driven workflows, troubleshoot production systems, and reduce manual effort through automation.

PROJECT EXPERIENCE

Multiverz - Challenges.one

Tools & Technologies - n8n, MongoDB, Redis, AWS S3, LLMs, Python, Node.js, REST APIs, Docker, LiteLLM, Web Search Tools

- Built an end-to-end multi-agent AI workflow covering domain analysis, rapid insights, foresight, consolidation, voting, capstone processing, and final report generation.
- Implemented parallel and sequential execution with checkpointing and retry/resume support to improve reliability of long-running AI workflows.
- Used LiteLLM for model routing and fallback handling across multiple LLMs, and integrated web search as a tool to enrich agent outputs with external context.
- Integrated MongoDB, Redis, AWS S3, REST APIs, and dynamic prompts to support workflow state, caching, document handling, and reusable client-specific automation.

TECHNICAL SKILLS

AI / ML Tools & Frameworks

LangGraph, LangChain, n8n, Agentic AI, Voice AI, RAG, Tool Calling, Prompt Engineering

LLMs & Model Platforms

OpenAI, Gemini, Claude, DeepSeek, Qwen, OpenRouter, LiteLLM

Programming Languages

Python, Java, JavaScript, TypeScript

Frontend Technologies

React, Tailwind CSS, Bootstrap

Backend Technologies

FastAPI, Spring Boot, Spring Security, JPA/Hibernate, REST APIs, Webhooks

Databases

PostgreSQL, MySQL, MongoDB, Redis, Qdrant, FAISS, Pinecone

Delivery and Observability

Docker, AWS EC2/S3, Git, SonarQube, Postman, Grafana, Langfuse

Testing

JUnit, Mockito, API testing with Postman

EDUCATION

Bachelor of Technology (Information Technology)

Conversational AI Voice Agent Platform

Tools & Technologies - Python, Pipecat, LangGraph, Gemini LLM, ElevenLabs, Twilio

- Built production-oriented multilingual AI voice agents for inbound and outbound conversations across Healthcare, BFSI, Insurance, and Retail, automating customer interactions and reducing manual call handling.
- Developed context-aware conversational agents using Python, Pipecat, LangGraph, Gemini, Twilio, and ElevenLabs, with tool calling for bookings, follow-ups, order tracking, loan processing, claims, and other business workflows.
- Implemented multilingual conversation flows with automatic language handling across English, Hindi, Gujarati, Marathi, and Spanish, enabling agents to adapt dynamically during customer interactions.
- Built reliable voice workflows with conversation state, retries, call handling, external API integrations, and reusable agent components, and containerized the solution with Docker for scalable deployment.

Cydra Social Multi-Agent SaaS

Tools & Technologies - Python, FastAPI, LangGraph, OpenAI/LLM APIs, React, PostgreSQL, Social APIs, OAuth, Webhooks

- Built and orchestrated specialized Viral, SEO, Community, and Curator agents within a multi-agent system for content generation, optimization, quality review, and platform-specific adaptation.
- Designed AI content workflows that generate multiple draft variations, refine outputs, and connect approved content to downstream scheduling and publishing actions.
- Integrated LLM-powered workflows with backend services, frontend interactions, and external social platforms through APIs, OAuth, and webhooks for end-to-end AI-driven automation.
- Improved AI workflow reliability through persistent state, asynchronous processing, failure handling, and fixes for concurrent content-generation issues.

WishAI - Internal RAG & AI Assistant

Tools & Technologies - React, Python (FastAPI), Qdrant, LangGraph, LangChain, PostgreSQL, LLM

- Built WishAI, an internal RAG-based AI assistant that helps employees quickly find organizational knowledge, documents, and resources without manually searching across scattered information sources.
- Implemented document ingestion, chunking, embeddings, metadata filtering, and Qdrant vector search to retrieve relevant context and generate accurate, grounded LLM responses.
- Added dynamic LLM selection and response caching to reduce repeated model calls, optimize token usage, and improve response performance.
- Implemented input/output guardrails, structured validation, and logging to support safer, more reliable, and controlled AI interactions.
- Added token monitoring with Grafana and integrated FastAPI-based LangGraph workflows with React chat interfaces for end-to-end conversational AI experiences.